

Iris Classification with Machine Learning

```
In [1]: #Importing Necessary Libraries
#Preprocessing Libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report, accuracy_score, f1_score

#Feature Selection Libraries
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression

# ML Libraries (Random Forest, Naive Bayes, SVM)
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC

# Evaluation Metrics
from yellowbrick.classifier import ClassificationReport
from sklearn import metrics

In [2]: #Read Iris Dataset into DataFrame & Print sample dataset
df = pd.read_csv('IrisClassification.csv')
df.head()
```

Out[2]:		Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1		5.1	3.5	1.4	0.2	Iris-setosa
1	2		4.9	3.0	1.4	0.2	Iris-setosa
2	3		4.7	3.2	1.3	0.2	Iris-setosa
3	4		4.6	3.1	1.5	0.2	Iris-setosa
4	5		5.0	3.6	1.4	0.2	Iris-setosa

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
# Column Non-Null Count Dtype
---  ---
0 Id 150 non-null int64
1 SepalLengthCm 150 non-null float64
2 SepalWidthCm 150 non-null float64
3 PetalLengthCm 150 non-null float64
4 PetalWidthCm 150 non-null float64
5 Species 150 non-null object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB

In [4]: df = df.drop(['Id'], axis=1)

In [5]: df['Species'].unique()

array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)

Out[5]:

In [6]: df['Species'] = pd.factorize(df['Species'])[0]
Target = 'Species'
df['Species'].unique()

Out[6]: array([0, 1, 2], dtype=int64)

In [7]: Features = ['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']
print('Full Features: ', Features)

Full Features: ['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']

In [8]: # Feature Selection using Recursive Feature Elimination
# Split DataFrame to target class and features
X_fs = df[Features]
Y_fs = df[Target]

# Feature Selection Model Fitting
model = LogisticRegression(solver='lbfgs', multi_class='auto')
```

```
In [9]: #Split dataset to Training Set & Test Set
x, y = train_test_split(df,
                        test_size = 0.2,
                        train_size = 0.8,
                        random_state= 3)

x1 = x[Features] #Features to train
x2 = x[Target] #Target Class to train
y1 = y[Features] #Features to test
y2 = y[Target] #Target Class to test

print('Feature Set Used : ', Features)
print('Target Class : ', Target)
print('Training Set Size : ', x.shape)
print('Test Set Size : ', y.shape)

Feature Set Used : ['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']
Target Class : Species
Training Set Size : (120, 5)
Test Set Size : (30, 5)

In [10]: # Gaussian Naive Bayes
# Create Model with configuration
nb_model = GaussianNB()

# Model Training
nb_model.fit(X=x1, y=y2)

# Prediction with Test Set
result= nb_model.predict(y[Features])
```

```
In [11]: # Model Evaluation
ac_sc = accuracy_score(y2, result)
rc_sc = recall_score(y2, result, average="weighted")
pr_sc = precision_score(y2, result, average="weighted")
f1_sc = f1_score(y2, result, average="micro")
confusion_m = confusion_matrix(y2, result)

print("===== Naive Bayes Results =====")
print("Accuracy : ", ac_sc)
print("Recall : ", rc_sc)
print("Precision : ", pr_sc)
print("F1 Score : ", f1_sc)
print("Confusion Matrix: ")
print(confusion_m)

===== Naive Bayes Results =====
Accuracy : 0.9666666666666667
Recall : 0.9666666666666667
Precision : 0.9696969696969696
F1 Score : 0.9666666666666667
Confusion Matrix:
[[10 0 0]
 [0 10 0]
 [0 1 9]]
```

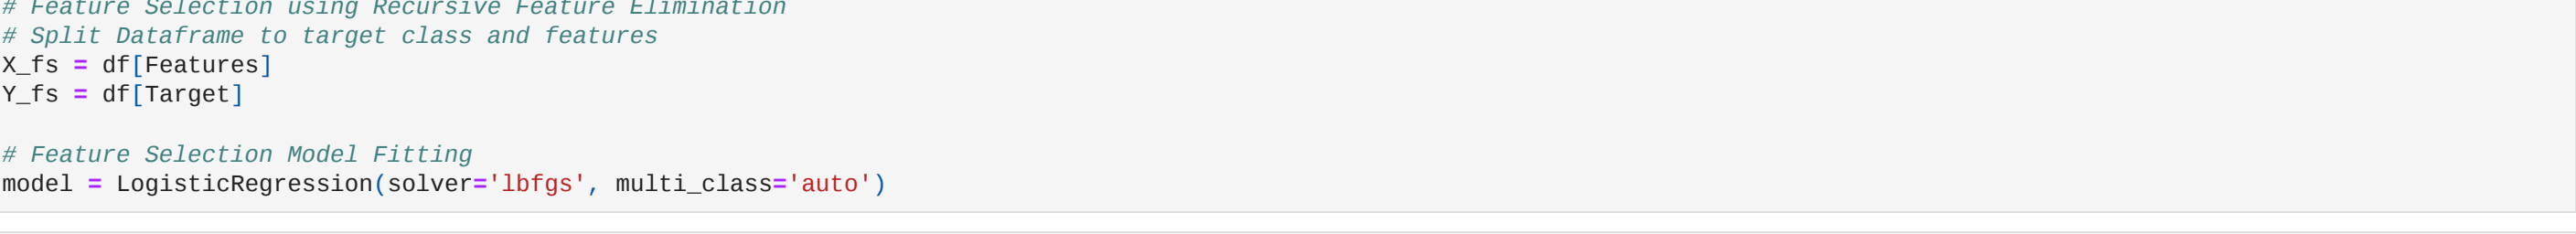
```
In [12]: # Classification Report
# Instantiate the classification model and visualizer
target_names = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
visualizer = ClassificationReport(nb_model, classes=target_names)
visualizer.fit(X=x1, y=y2) # Fit the training data to the visualizer
visualizer.score(y1, y2) # Evaluate the model on the test data

print('===== Classification Report =====')
print('')
print(classification_report(y2, result, target_names=target_names))

g = visualizer.poof()

===== Classification Report =====
```

C:\Users\ushaj\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but GaussianNB was fitted with feature names



```
In [13]: # Random Forest
# Create Model with configuration
rf_model = RandomForestClassifier(n_estimators=70, # Number of trees
                                min_samples_split = 30,
                                bootstrap = True,
                                max_depth = 50,
                                min_samples_leaf = 25)

# Model Training
rf_model.fit(X=x1, y=y2)

# Prediction
result = rf_model.predict(y[Features])
```

```
In [14]: # Model Evaluation
ac_sc = accuracy_score(y2, result)
rc_sc = recall_score(y2, result, average="weighted")
pr_sc = precision_score(y2, result, average="weighted")
f1_sc = f1_score(y2, result, average="micro")
confusion_m = confusion_matrix(y2, result)

print("===== Random Forest Results =====")
print("Accuracy : ", ac_sc)
print("Recall : ", rc_sc)
print("Precision : ", pr_sc)
print("F1 Score : ", f1_sc)
print("Confusion Matrix: ")
print(confusion_m)

===== Random Forest Results =====
Accuracy : 0.9666666666666667
Recall : 0.9666666666666667
Precision : 0.9696969696969696
F1 Score : 0.9666666666666667
Confusion Matrix:
[[10 0 0]
 [0 10 0]
 [0 1 9]]
```


```
In [15]: # Classification Report
# Instantiate the classification model and visualizer
target_names = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
visualizer = ClassificationReport(rf_model, classes=target_names)
visualizer.fit(X=x1, y=y2) # Fit the training data to the visualizer
visualizer.score(y1, y2) # Evaluate the model on the test data

print('===== Classification Report =====')
print('')
print(classification_report(y2, result, target_names=target_names))

g = visualizer.poof()

===== Classification Report =====
```

C:\Users\ushaj\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names



```
In [16]: # Support Vector Machine
# Create Model with configuration
svm_model = SVC(kernel='linear')

# Model Training
svm_model.fit(X=x1, y=y2)

# Prediction
result = svm_model.predict(y[Features])
```

```
In [17]: # Model Evaluation
ac_sc = accuracy_score(y2, result)
rc_sc = recall_score(y2, result, average="weighted")
pr_sc = precision_score(y2, result, average="weighted")
f1_sc = f1_score(y2, result, average="micro")
confusion_m = confusion_matrix(y2, result)

print("===== SVM Results =====")
print("Accuracy : ", ac_sc)
print("Recall : ", rc_sc)
print("Precision : ", pr_sc)
print("F1 Score : ", f1_sc)
print("Confusion Matrix: ")
print(confusion_m)

===== SVM Results =====
Accuracy : 0.9666666666666667
Recall : 0.9666666666666667
Precision : 0.9696969696969696
F1 Score : 0.9666666666666667
Confusion Matrix:
[[10 0 0]
 [0 9 1]
 [0 0 10]]
```

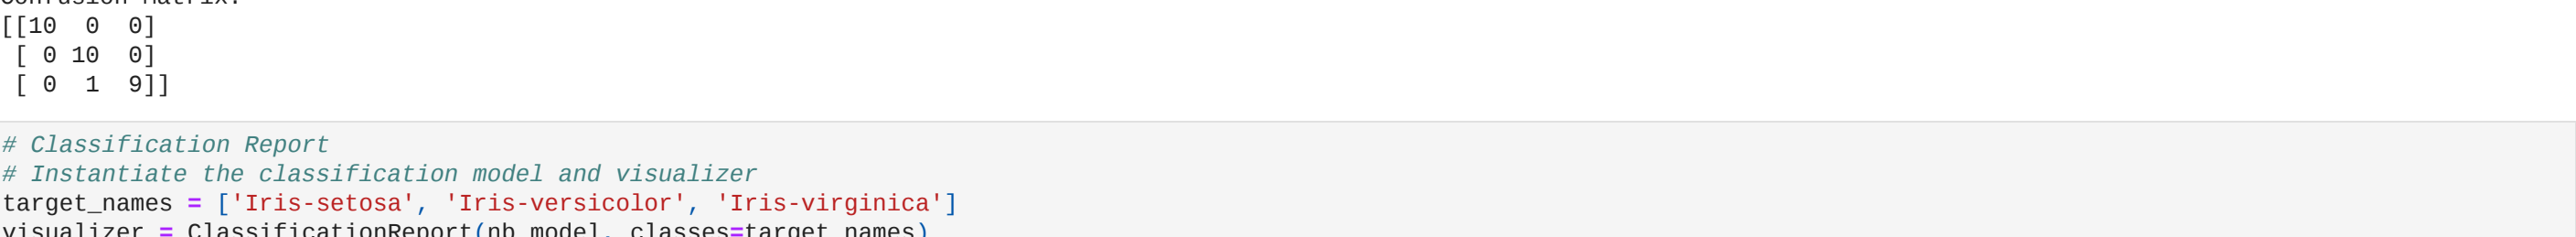
```
In [18]: # Classification Report
# Instantiate the classification model and visualizer
target_names = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
visualizer = ClassificationReport(svm_model, classes=target_names)
visualizer.fit(X=x1, y=y2) # Fit the training data to the visualizer
visualizer.score(y1, y2) # Evaluate the model on the Test Set

print('===== Classification Report =====')
print('')
print(classification_report(y2, result, target_names=target_names))

g = visualizer.poof()

===== Classification Report =====
```

C:\Users\ushaj\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but SVC was fitted with feature names



```
In [19]: # Ensemble Voting Model
# Combine 3 Models to create an Ensemble Model

# Create Model with configuration
ecclf1 = VotingClassifier(estimators=[('svm', svm_model), ('rf', rf_model), ('gnb', nb_model)],
                          weights=[1,1,1],
                          flatten_transform=True)

ecclf1 = ecclf1.fit(X=x1, y=y2)

# Prediction
result = ecclf1.predict(y[Features])
```

```
In [20]: # Model Evaluation
ac_sc = accuracy_score(y2, result)
rc_sc = recall_score(y2, result, average="weighted")
pr_sc = precision_score(y2, result, average="weighted")
f1_sc = f1_score(y2, result, average="micro")
confusion_m = confusion_matrix(y2, result)

print("===== Ensemble Voting Results =====")
print("Accuracy : ", ac_sc)
print("Recall : ", rc_sc)
print("Precision : ", pr_sc)
print("F1 Score : ", f1_sc)
print("Confusion Matrix: ")
print(confusion_m)

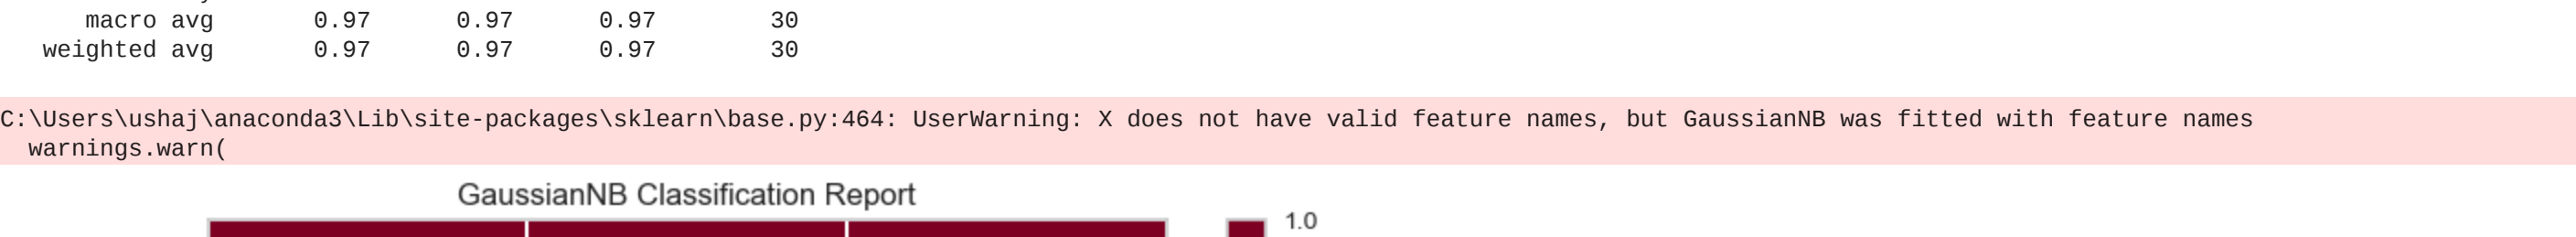
===== Ensemble Voting Results =====
Accuracy : 0.9666666666666667
Recall : 0.9666666666666667
Precision : 0.9696969696969696
F1 Score : 0.9666666666666667
Confusion Matrix:
[[10 0 0]
 [0 10 0]
 [0 1 9]]
```

```
In [21]: # Classification Report
# Instantiate the classification model and visualizer
target_names = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
visualizer = ClassificationReport(ecclf1, classes=target_names)
visualizer.fit(X=x1, y=y2) # Fit the training data to the visualizer
visualizer.score(y1, y2) # Evaluate the model on the test data

print('===== Classification Report =====')
print('')
print(classification_report(y2, result, target_names=target_names))

g = visualizer.poof()

===== Classification Report =====
```



C:\Users\ushaj\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but SVC was fitted with feature names

